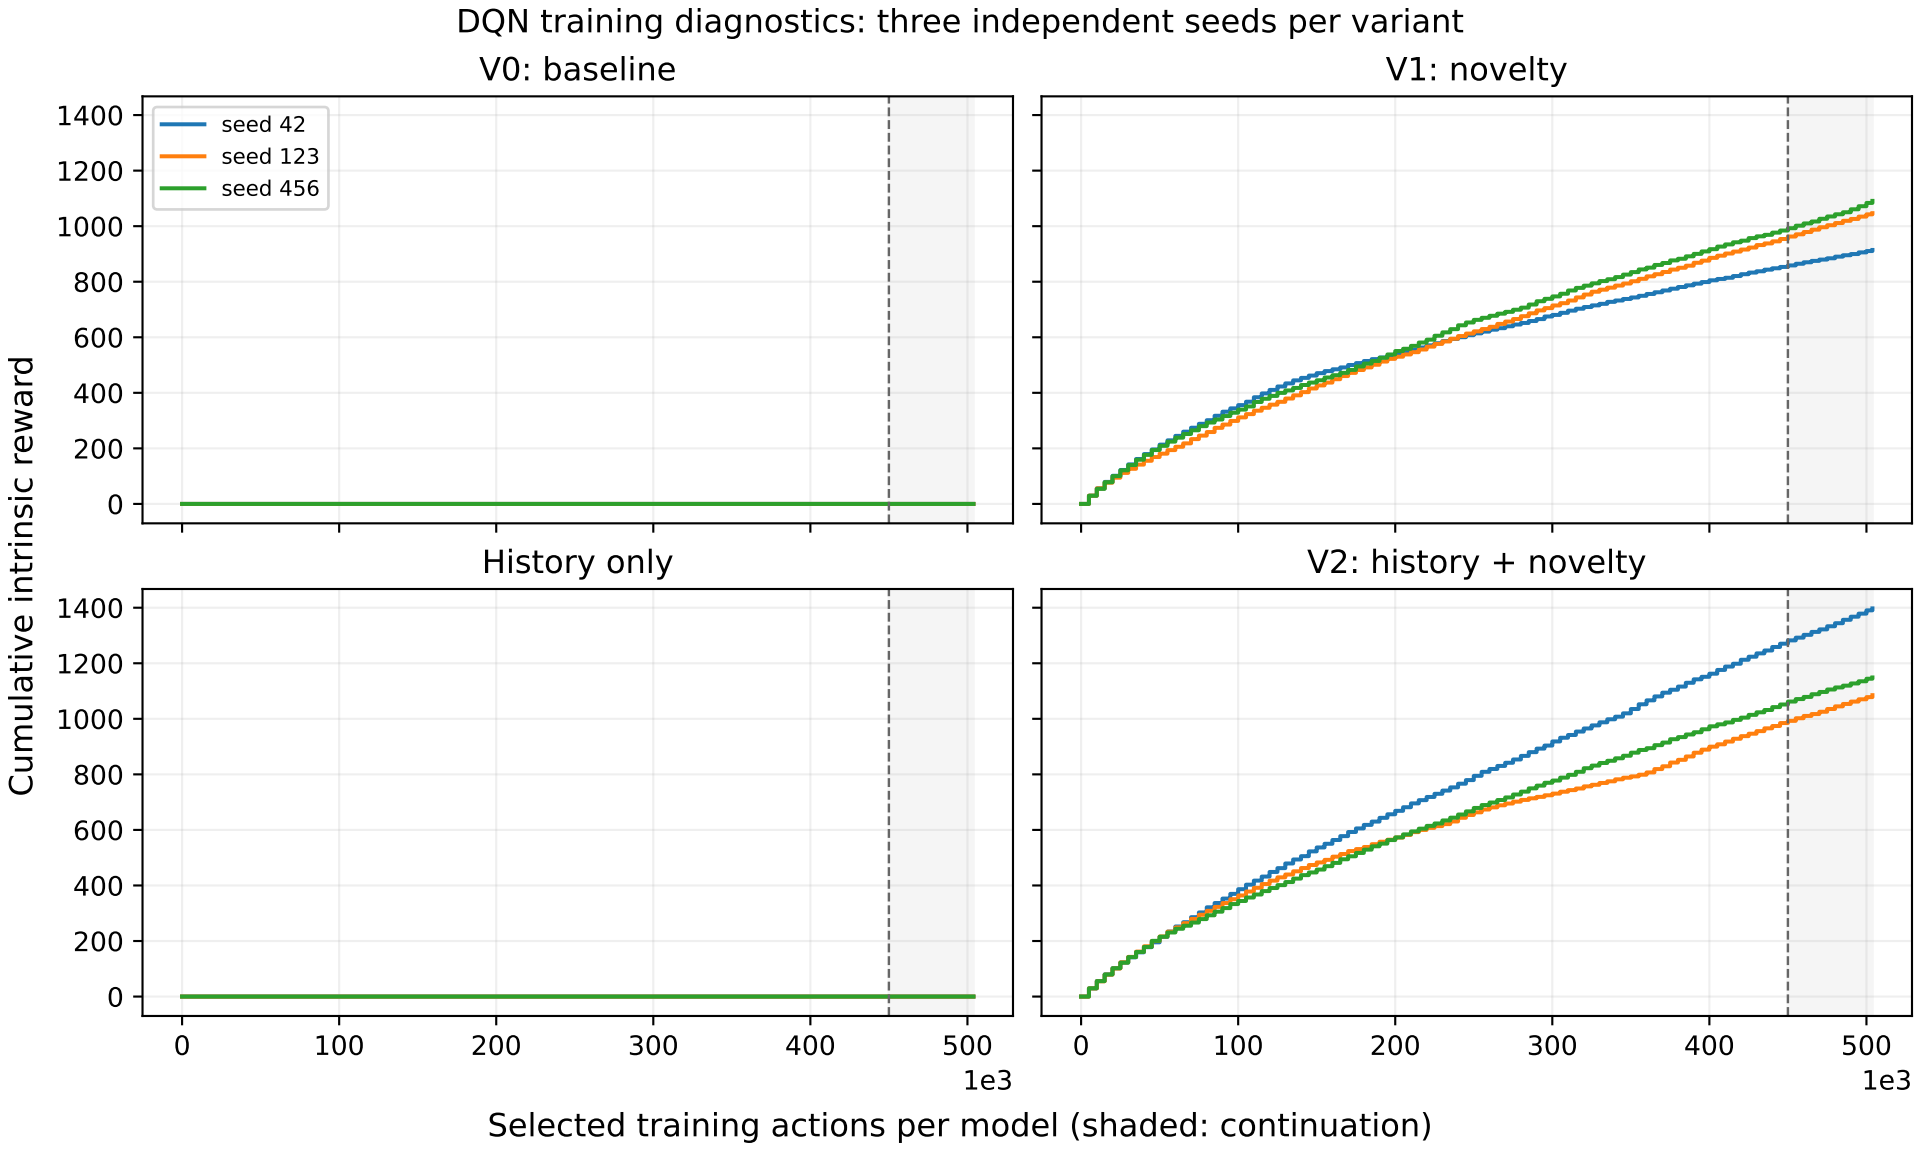
DQN training diagnostics: three independent seeds per variant
V0: baseline
V1: novelty
History only
V2: history + novelty
seed 42
seed 123
seed 456
Cumulative intrinsic reward
0
200
400
600
800
1000
1200
1400
0
100
200
300
400
500
1e3
Selected training actions per model (shaded: continuation)