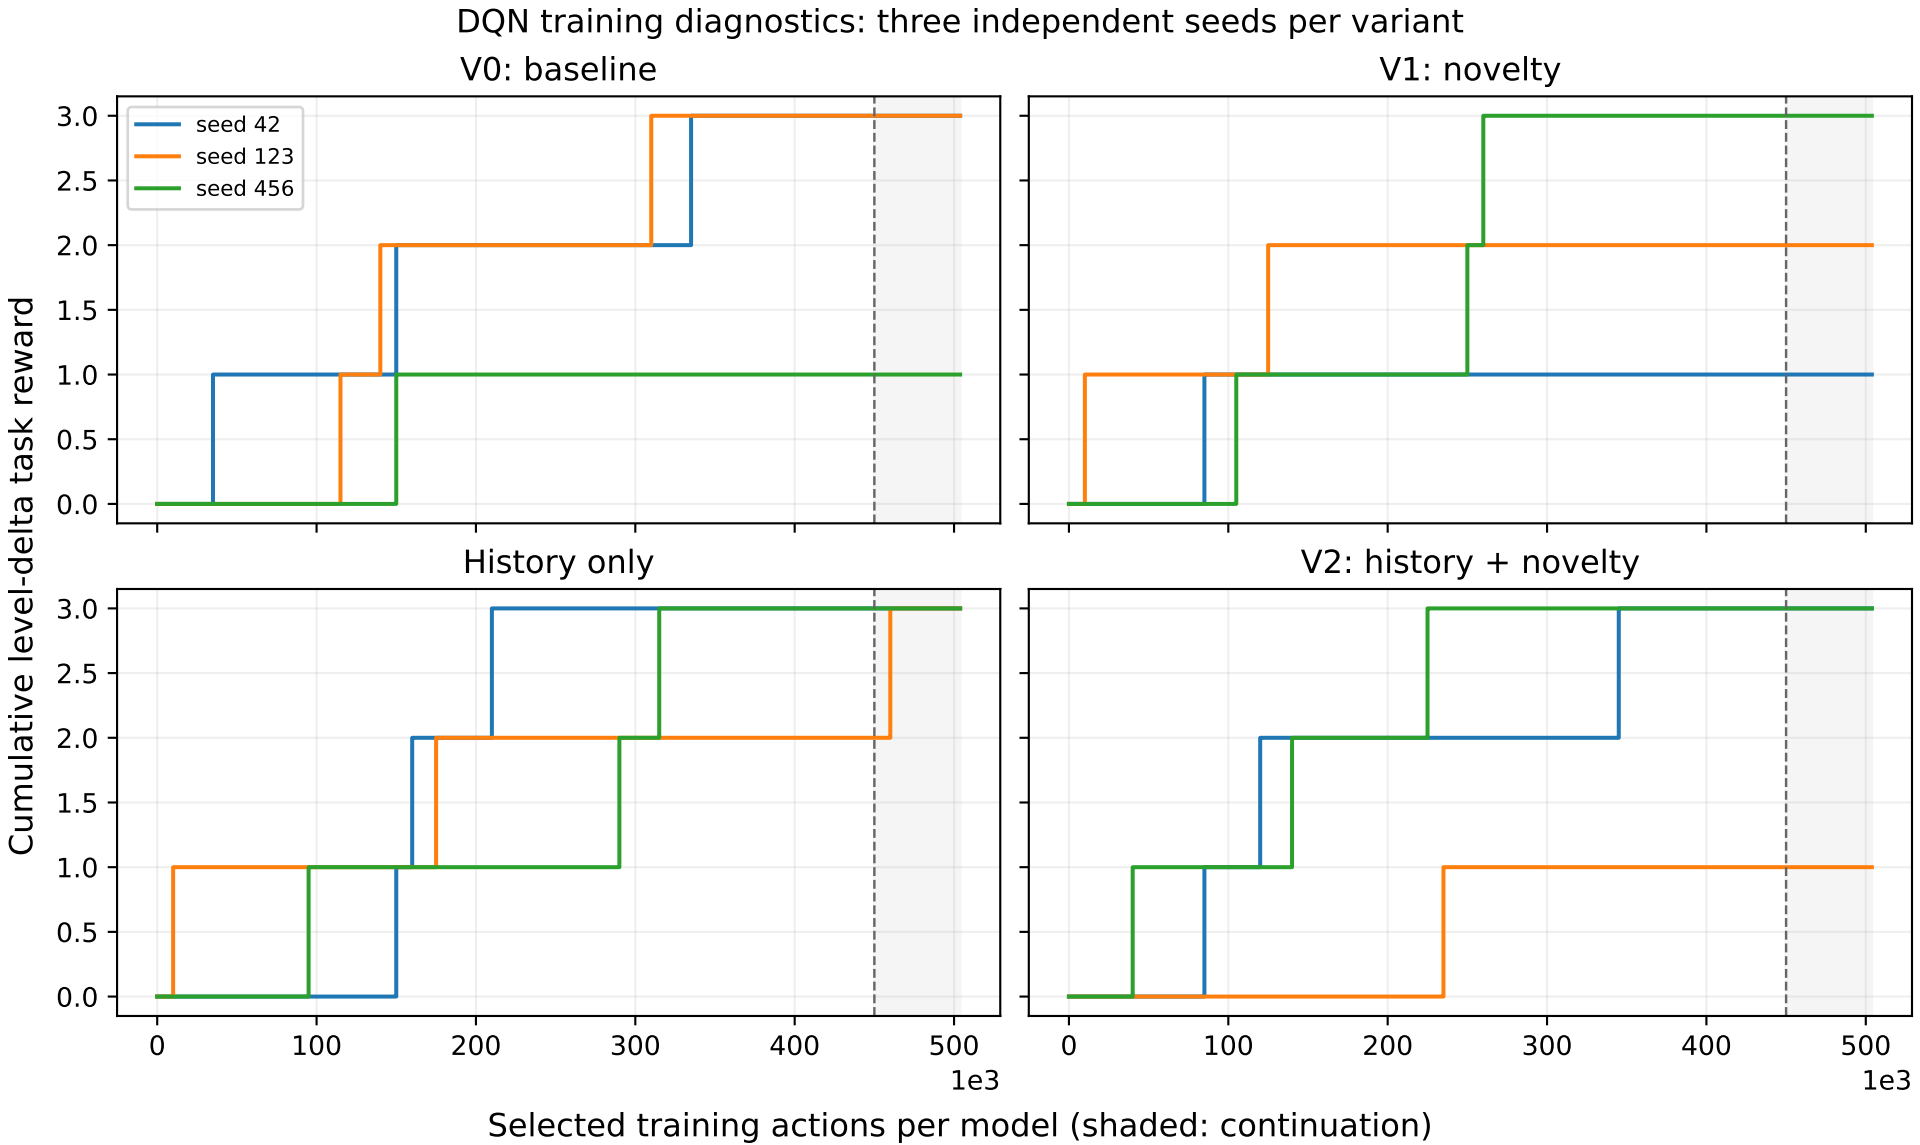
DQN training diagnostics: three independent seeds per variant
V0: baseline
V1: novelty
History only
V2: history + novelty
seed 42
seed 123
seed 456
Cumulative level-delta task reward
Selected training actions per model (shaded: continuation)
1e3
0.0
0.5
1.0
1.5
2.0
2.5
3.0
0
100
200
300
400
500